

# Intel Compiler Options and Optimizations

SNUG TechTalk

SciNet

[www.scinet.utoronto.ca](http://www.scinet.utoronto.ca)

University of Toronto

Toronto, Canada

December 14, 2011

# Outline

- 1 Compiler
- 2 Optimization
- 3 Other Options
- 4 MKL

## Intel® Composer XE 2011 aka (v12.1)

- C (icc)
- C++ (icpc)
- FORTRAN (ifort)
- Threaded Building Blocks (TBB)
- Integrated Performance Primitives (IPP)
- Math Kernel Libraries (MKL)

## Optimization Levels

- **-O0** disable optimization
- **-O1** optimizes for code size
- **-O2** optimizes for speed (default)
- **-O3** **-O2** plus more aggressive optimizations

## Optimization Levels

- **-O0** disable optimization
- **-O1** optimizes for code size
- **-O2** optimizes for speed (default)
- **-O3** **-O2** plus more aggressive optimizations

## From the Manual

“The **-O3** option is particularly recommended for applications that have loops that do many floating-point calculations or process large data sets.”

## -02 Optimizations

- intrinsic inlining
- inlining
- constant propagation
- forward substitution
- routine attribute propagation
- variable address-taken analysis
- dead static function elimination
- removal of unreferenced variables
- constant propagation
- copy propagation
- dead-code elimination
- global register allocation
- global instruction scheduling and control speculation
- loop unrolling
- optimized code selection
- partial redundancy elimination
- strength reduction/induction variable simplification
- variable renaming
- exception handling optimizations
- tail recursions
- peephole optimizations
- structure assignment lowering and optimizations
- dead store elimination

# Optimization Terminology

## Inlining

### Inlining

Replaces the function call with the actual functions code.

# Optimization Terminology

## Inlining

Replaces the function call with the actual functions code.

## Original

```
int func(int &x,int &y) { return 4*x+3*y; }

int main(){
    int x=4, y=3;
    int b=fun(x,y)
}
```



# Optimization Terminology

## Inlining

Replaces the function call with the actual functions code.

## Original

```
int func(int &x,int &y) { return 4*x+3*y; }

int main(){
    int x=4, y=3;
    int b=fun(x,y)
}
```

## Inlined

```
int main(){
    int x=4,y=3;
    int b= 4*x+3*y;
}
```

# Optimization Terminology

## Branch Elimination

### Original

```
if ( x < x1 ) {  
    a = a0 + a1;  
} else if ( x < x2 ) {  
    a = a0 - a1;  
} else if ( x < x3 ) {  
    a = a0 * a1;  
} else if ( x < x4 ) {  
    a = a0 / a1;  
} else {  
    a = a0;  
}
```

### Optimizer Approaches

- static branch elimination
- compute all cases and conditions, then pick the correct one
- replace with switch statements, jump tables
- branch re-alignment

et

## -O3 Additional Optimizations

- Loop Blocking for cache
- Loop Permutation or Interchange
- Loop Distribution
- Loop Fusion
- Loop Unrolling
- Unroll and Jam
- Loop Blocking or Tiling
- Loop Reversal
- Loop Peeling
- Loop Rerolling
- Profile-Guided Loop Unrolling
- Code Replication to eliminate branches
- Memory-access optimizations
- Data Prefetching
- Scalar Replacement
- Partial-Sum Optimization
- Predicate Optimization
- Data Transformation: Malloc Combining and Memset Combining
- Memset and Malloc Recognition
- Statement Sinking for Creating Perfect Loopnests

# Optimization Terminology

## Loop Unrolling

### Original

```
for (int x=0; x < 100; x++)  
{  
    func(x);  
}
```

# Optimization Terminology

## Loop Unrolling

### Original

```
for (int x=0; x < 100; x++)  
{  
    func(x);  
}
```

### Optimized

```
for (int x = 0; x < 100; x+=5)  
{  
    func(x);  
    func(x+1);  
    func(x+2);  
    func(x+3);  
    func(x+4);  
}
```

# Optimization Terminology

## Loop Collapsing

### Original

```
int a[100][300];  
for (int i = 0; i < 300; i++)  
    for (int j = 0; j < 100; j++)  
        a[j][i] = 0;
```

# Optimization Terminology

## Loop Collapsing

### Original

```
int a[100][300];  
for (int i = 0; i < 300; i++)  
    for (int j = 0; j < 100; j++)  
        a[j][i] = 0;
```

### Optimized

```
int a[100][300];  
int *p = &a[0][0];  
  
for (int i = 0; i < 30000; i++)  
    *p++ = 0;
```

# Optimization Terminology

## Loop Fusion

### Original

```
int x[100], y[100];  
for (int i = 0; i < 100; i++)  
    x[i] = 1;  
for (int i = 0; i < 100; i++)  
    y[i] = 2;
```



# Optimization Terminology

## Loop Fusion

### Original

```
int x[100], y[100];  
for (int i = 0; i < 100; i++)  
    x[i] = 1;  
for (int i = 0; i < 100; i++)  
    y[i] = 2;
```

### Optimized

```
int x[100], y[100];  
for (int i = 0; i < 100; i++)  
{  
    x[i] = 1;  
    y[i] = 2;  
}
```

# Optimization Terminology

## Loop Peeling

### Original

```
int p = 10;
for (int i=0; i<10; ++i)
{
    y[i] = x[i] + x[p];
    p = i;
}
```

# Optimization Terminology

## Loop Peeling

### Original

```
int p = 10;
for (int i=0; i<10; ++i)
{
    y[i] = x[i] + x[p];
    p = i;
}
```

### Optimized

```
y[0] = x[0] + x[10];
for (int i=1; i<10; ++i)
{
    y[i] = x[i] + x[i-1];
}
```

## System Specific

- **-march="cpu"** optimize for a specific cpu
- **-mtune="cpu"** produce code only for a specific cpu
- **-msse3,-msse4,-mavx, etc.** level of SIMD and vector instructions

## System Specific

- **-march="cpu"** optimize for a specific cpu
- **-mtune="cpu"** produce code only for a specific cpu
- **-msse3,-msse4,-mavx, etc.** level of SIMD and vector instructions

## Use this instead!

**-xHost** optimize and tune for the compiling CPU

# Optimizations

## System Specific

- **-march="cpu"** optimize for a specific cpu
- **-mtune="cpu"** produce code only for a specific cpu
- **-msse3,-msse4,-mavx, etc.** level of SIMD and vector instructions

## Use this instead!

**-xHost** optimize and tune for the compiling CPU

## GPC Recommendations

**-xHost -O3**

# Optimization Terminology

## Vector Extensions

### Intel x86\_64 extensions

- Streaming SIMD Extensions (SSE1 - SSE4.2)
- AVX

### Original x86

Add two single precision vectors requires four floating-point addition instructions.

```
vec_res.x = v1.x + v2.x;  
vec_res.y = v1.y + v2.y;  
vec_res.z = v1.z + v2.z;  
vec_res.w = v1.w + v2.w;
```

# Optimization Terminology

## Vector Extensions

### Intel x86\_64 extensions

- Streaming SIMD Extensions (SSE1 - SSE4.2)
- AVX

### Original x86

Add two single precision vectors requires four floating-point addition instructions.

```
vec_res.x = v1.x + v2.x;  
vec_res.y = v1.y + v2.y;  
vec_res.z = v1.z + v2.z;  
vec_res.w = v1.w + v2.w;
```

### SSE

A single 128-bit 'packed-add' replaces four scalar addition instructions.

```
movaps xmm0, [v1]; xmm0 = v1.w | v1.z | v1.y | v1.x  
addps xmm0, [v2];  xmm0 = v1.w+v2.w | v1.z+v2.z | v1.y+v2.y | v1.x+v2.x  
movaps [vec_res], xmm0
```



## Inter Procedural Optimizations (IPO)

- inlining
- constant propagation
- mod/ref analysis
- alias analysis
- forward substitution
- routine key-attribute propagation
- address-taken analysis
- partial dead call elimination
- symbol table data promotion
- common block variable coalescing
- dead function elimination
- unreferenced variable removal
- whole program analysis
- array dimension padding
- common block splitting
- stack frame alignment
- structure splitting and field reordering
- formal parameter alignment analysis
- C++ class hierarchy analysis
- indirect call conversion
- specialization
- Passing arguments in registers to optimize calls and register usage

## Inter Procedural Optimizations

- **-ip** single file ip optimization
- **-ipo** multiple file or whole program optimization

## Inter Procedural Optimizations

- **-ip** single file ip optimization
- **-ipo** multiple file or whole program optimization

## Profile Guided Optimizations

- **-prof-gen** instrument code to generate profile
- **-prof-use** use profile to guide optimization

# Optimizations

## Inter Procedural Optimizations

- **-ip** single file ip optimization
- **-ipo** multiple file or whole program optimization

## Profile Guided Optimizations

- **-prof-gen** instrument code to generate profile
- **-prof-use** use profile to guide optimization

## Flank Speed

**-fast** enables **-xHost -O3 -ipo -no-prec-div -static**

## -fpmode

- **fast=1** default
- **fast=2** most aggressive
- **precise** value-safe optimizations on intermediate operations
- **except** strict floating point semantics
- **strict** disables all “fast-math” options

## If Required

For floating point consistency and reproducibility use:

**-fpmode precise -fpmode except**

Seen this error?

relocation truncated to fit: R\_X86\_64\_PC32

Seen this error?

relocation truncated to fit: R\_X86\_64\_PC32

**-mcmmodel=**

- **small** code and data restricted to the first 2GB of address space
- **medium** code restricted to the first 2GB of address space
- **large** no restrictions

## MKL Components

- BLAS
- LAPACK
- ScaLAPACK
- FFT
- PBLAS
- BLACS
- plus others



## Link Line - MKL 10.3 or less

**-lmkl\_intel\_lp64 -lmkl\_sequential -lmkl\_core -lpthread -lm**

## Link Line - Composer XE 2011

- **-mkl=sequential** no-threaded versions (serial)
- **-mkl=parallel** threaded (openmp)
- **-mkl=cluster** for ScaLAPACK, FFT, BLACS

## Link Line Advisor

<http://software.intel.com/en-us/articles/intel-mkl-link-line-advisor/>

## Intel Documentation

<http://software.intel.com/en-us/articles/intel-parallel-studio-xe-for-linux-documentation/>

## Compiler Optimization flags

[http://software.intel.com/sites/products/collateral/hpc/compilers/compiler\\_qrg12.pdf](http://software.intel.com/sites/products/collateral/hpc/compilers/compiler_qrg12.pdf)

## White Paper on Floating Point

[https://support.scinet.utoronto.ca/wiki/images/f/f2/FP\\_Consistency.pdf](https://support.scinet.utoronto.ca/wiki/images/f/f2/FP_Consistency.pdf)